

Blockchains & Cryptocurrencies

Alternative Puzzles - II



Instructor: Abhishek Jain
Johns Hopkins University - Spring 2021

*Some slides based on NBFMG

Housekeeping

- HW1 Due Today
- Course Project Idea Due on 02/26

Last Time

- Mining Strategies (Default and Non-Default) & Incentives
- Puzzle Requirements and Alternative Puzzles (Motivation and Designs)

Today

- Alternative Puzzles (contd.)

Recall: Puzzle requirements

- Cheap to Verify
 - since other users have to verify solutions
- Adjustable difficulty
 - E.g., due on increase in hash rate or more users
- In PoW puzzles, chance of winning should be proportional to computing power (e.g., hash power in Bitcoin)
 - Large players get only proportional advantage
 - Even small players get proportional compensation

ASIC Resistant (PoW) Puzzles

ASIC resistance - Why? (1 of 2)

Goal: Ordinary people with idle laptops, PCs, or even mobile phones can mine!

Lower barrier to entry

Approach: Reduce the gap between custom hardware and general purpose equipment

ASIC resistance - Why? (2 of 2)

Goal: Prevent large manufacturers from dominating the game

“Burn-in” advantage


In-house designs



Approach: reduce the “gap” between future hardware and the custom ASICs we already have

Memory hard puzzles

Premise: the cost and performance of memory is more stable than for processors



script

Colin Percival, 2009

- Memory hard hash function
 - ***Constant time/memory tradeoff***
 - Memory consumes a large amount of on-chip area. High memory requirement => small number of hashing engines on special-purpose chips
- Widely used alternative PoW puzzle (e.g., Litecoin)
- Also used in Password-hashing

1. Fill memory with random values

2. Read from the memory in random order

script - step 1 of 2 (write)

Input: \mathbf{x}


$$\mathbf{V}_1 = \mathbf{H}(\mathbf{x})$$

$$\mathbf{V}_2 = \mathbf{H}(\mathbf{V}_1) = \mathbf{H}(\mathbf{H}(\mathbf{x}))$$

$$\mathbf{V}_3 = \mathbf{H}(\mathbf{V}_2) = \mathbf{H}^3(\mathbf{x})$$

...

$$\mathbf{V}_N = \mathbf{H}^N(\mathbf{x})$$



script - step 2 of 2 (read)

Input: \mathbf{x}


$\mathbf{A} := \mathbf{H}^{N+1}(\mathbf{x})$

For N iterations:

$i := \mathbf{A} \bmod N$

$\mathbf{A} := \mathbf{H}(\mathbf{A} \text{ xor } \mathbf{V}_i)$


Output: \mathbf{A}



script - time/memory tradeoff

Why is this memory-hard?

Reduce memory by half, 1.5x the # steps



Need to access V_i where i is even?

Access V_{i-1}

Compute $V_i = H(V_{i-1})$

script

Disadvantages: Also requires N steps, N memory to check

Is it actually ASIC resistant?

script ASICs **are** already available

Exploit time-memory trade-offs, lower values of N , etc.

Academic research

- Many subsequent candidates: Argon2i (winner of PW-hashing contest), Ballon-Hashing, etc.
- Proofs of memory hardness in various models using graph pebbling techniques (see, e.g., Alwen-Serbeninko'15 and many subsequent works)

Cuckoo hash cycles

John Tromp, 2014

(Conjectured) Memory hard puzzle that's cheap to verify

Input: X


For $i = 1$ to E :

$a := H_0(X + i)$

$b := N + H_1(X + i)$

edge($a \bmod N, b \bmod N$)

Is there a cycle of size K ? If so, Output: X, K edges



Even more approaches

- More complicated hash functions
 - XII: 11 different hash functions combined (subsequent iterations: XI3, XI4, XI5, XI7)
- Moving target

Change the puzzle periodically

Counter argument: SHA2 is fine


Bitcoin Mining ASICs aren't changing much

Big ASICs only marginally more performant than small ones


Ordinary SHA2 Circuit



Affordable ASIC



Expensive ASIC



Proof-of-useful-work

Recovering wasted work

Recall: power consumed by Bitcoin network in 2019
~ power consumed by Switzerland :(

Natural question:

Can we recycle this and do something useful?

Candidates - needle in a haystack

- Potential choices:
 - Protein folding (find a low energy configuration)
 - Search for aliens (find an anomalous region of a signal)

Candidates - needle in a haystack

- Potential choices:
 - Protein folding (find a low energy configuration)
 - Search for aliens (find an anomalous region of a signal)
- Challenges:
 - Randomly chosen instances must be hard
 - Who chooses the problem?
 - Verification must also be efficient

Primecoin

Sunny King, 2013



Puzzle based on finding large prime numbers

Puzzle parameters: **m, n, k**

Cunningham chain:

p_1, p_2, \dots, p_k where $p_i = 2p_{i-1} + 1$

Each p_i is a large (probable) prime

p_1 is **n**-bit long and shares same leading **m** bits as

$H(\text{prev} \parallel \text{mrkl_root} \parallel \text{nonce})$

Primecoin



- Many of the largest known Cunningham chains have come from Primecoin miners
- Hard problem? Studied by others (e.g., PrimeGrid)
- Usefulness? Some applications to crypto (e.g., Young-Yung'98)

Short Paper: The Proof is in the Pudding

Proofs of Work for Solving Discrete Logarithms

Marcella Hastings¹, Nadia Heninger², and Eric Wustrow³

¹ University of Pennsylvania

² University of California, San Diego

³ University of Colorado Boulder

Abstract. We propose a proof of work protocol that computes the discrete logarithm of an element in a cyclic group. Individual provers generating proofs of work perform a distributed version of the Pollard rho algorithm. Such a protocol could capture the computational power expended to construct proof-of-work-based blockchains for a more useful purpose, as well as incentivize advances in hardware, software, or algorithms for an important cryptographic problem. We describe our proposed construction and elaborate on challenges and potential trade-offs that arise in designing a practical proof of work.

Keywords: Proofs of work, discrete log, Pollard rho

Recovering wasted hardware

Estimate: more than \$100M spent on customized Bitcoin mining hardware (old data from 2014)

This hardware investment is otherwise useless

Idea: a puzzle where hardware investment is useful, even if the work is wasted?

Permacoin - Mining with storage

Miller et al., 2014

Bitcoin



Permacoin



Side effect:

Massively distributed, replicated storage system

Permacoin


Assume we have a large file **F** to store

For simplicity: **F** is chosen globally, at the beginning, by a trusted dealer

Each user stores a random subset of the file

Storage-based puzzle

1. Build a Merkle tree, where each leaf is a segment of the file
2. Generate a public signing key pk , which determines a random subset of file segments to store
3. Each mining attempt:
 - a) Select a random nonce
 - b) $h1 := H(\text{prev} \parallel \text{mrkl_root} \parallel PK \parallel \text{nonce})$
 - c) $h1$ selects k segments from subset
 - d) $h2 := H(\text{prev} \parallel \text{mrkl_root} \parallel PK \parallel \text{nonce} \parallel F)$
 - e) Winner if $h2 < \text{TARGET}$



Summary

- Useful proof-of-work is a natural goal
(while maintaining security requirements)
- The benefit must be a pure public good
- Viable approaches include storage, prime-finding, others may be possible
- Realized benefit so far has been limited


Non-outsourceable Puzzles

Large mining pools are a threat

- Bitcoin's core value is decentralization
- If power is consolidated in a few large pools, the operators are targets for coercion/hacking
- Position: large pools should be discouraged
Analogy to voting: It's illegal (in US) to sell your vote

June 12, 2014

GHash.IO large mining pool crisis




Observation:

Pool participants don't trust each other

Pools only work because the “shares” protocol lets members ***prove*** cooperation

Standard Bitcoin mining pool



The Vigilante Attack

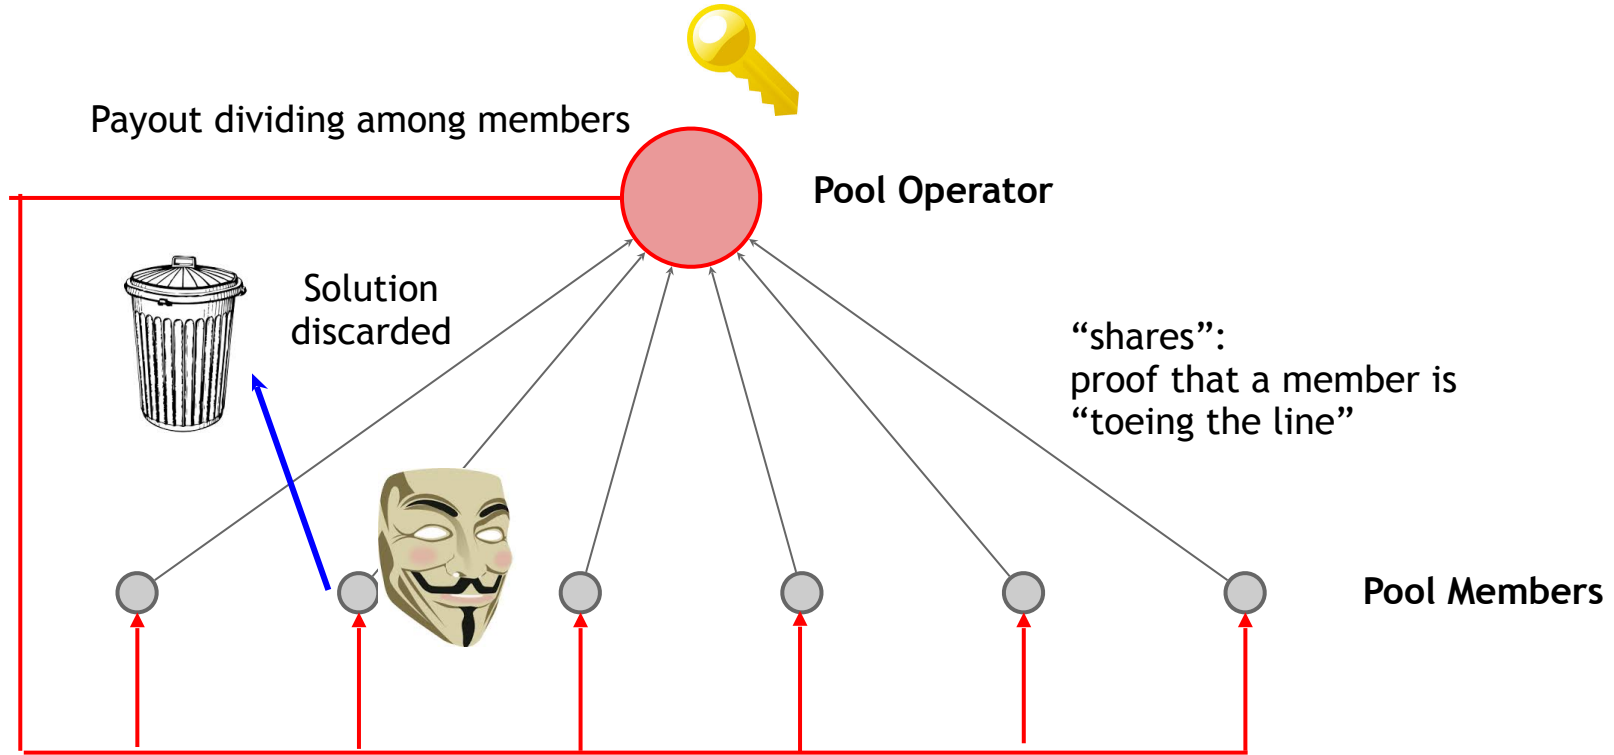
Suppose a Vigilante is angry with a large pool

He submits “shares” like normal....

... but if he finds a real solution, discards it

Pool output is reduced, Vigilante loses a little

The Vigilante Attack



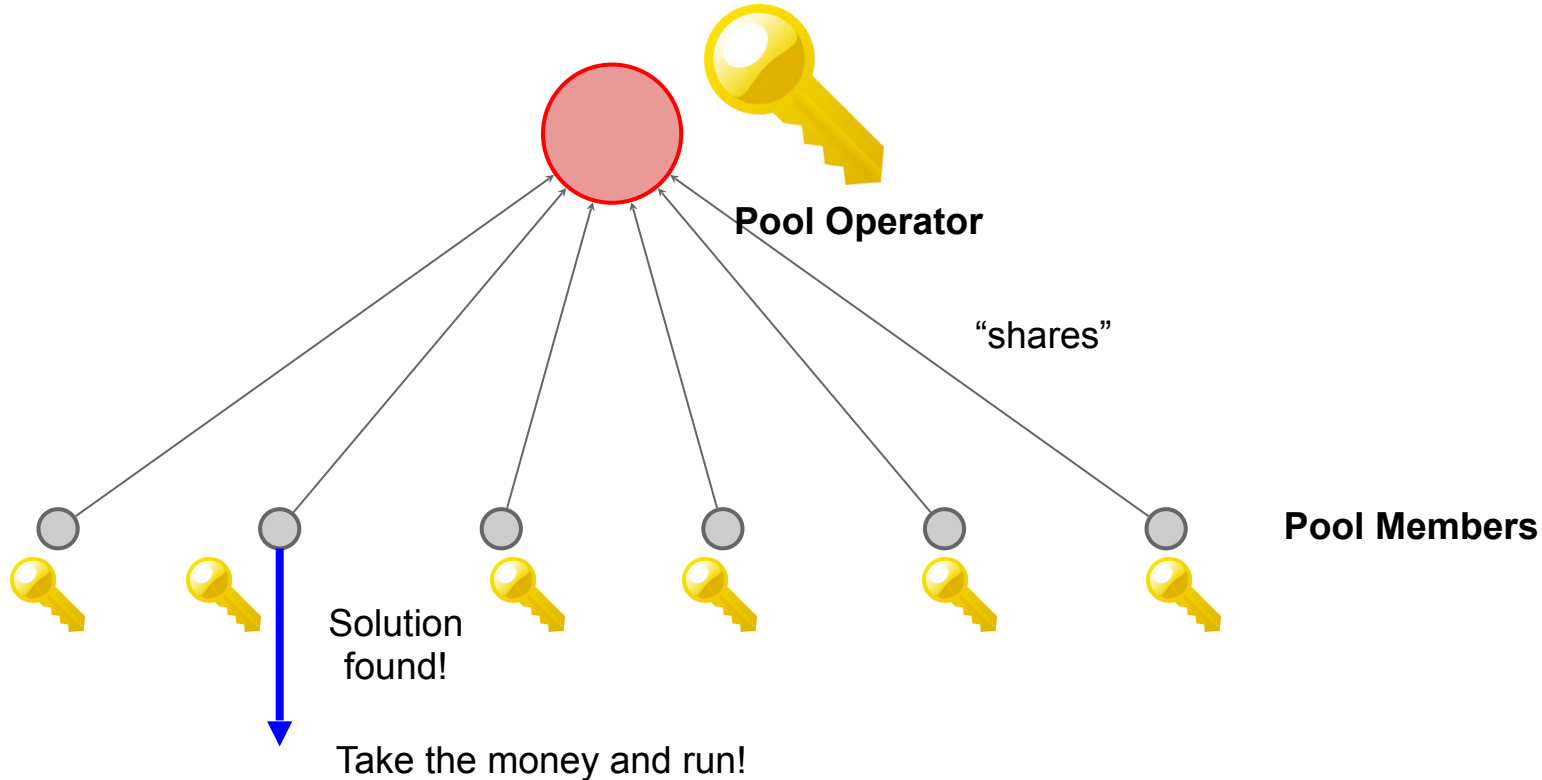
Encouraging the Vigilante

Whoever ***FINDS*** a solution spends the reward

Approach:

- searching for a solution requires ***SIGNING***, not just hashing. (Knowledge of a private key)
- Private key can be used to spend the reward

Encouraging the Vigilante



Nonoutsourcable puzzle

Solution:

`(prev, mrkl_root, nonce, PK, s1, s2)`

such that:

`H(prev || PK || nonce || s1) < TARGET`
`VerifySig(PK, s1, prev || nonce)`
`VerifySig(PK, s2, prev || mrkl_root)`

Signature needed to find solution

Public Key

s1


s2

Second signature spends reward

“Virtual Mining”

Bitcoin Mining has an unnecessary step

Proof-of-Work Mining:



Proof of Stake

- Creator of next block chosen at random based on amount of current “stake” in the system
- Assuming all the money owned/used by miners is in the system, this mechanism cuts the middle man (equipment manufacturer)

Potential benefits

- Lower overall costs
 - No harm to the environment
 - Savings distributed to all coin holders
- Stakeholder incentives - good stewards?
- No ASIC advantage
- 51% attack might be harder (this is debatable)

Examples of PoS based Cryptocurrencies

- Cardano
- Algorand
- Ethereum 2 (hopefully!)
- Nxt
- Neucoin
- ...

Examples of secure PoS systems

- Algorand [Full version: Chen-Micali'17]
- Cardano/Ourboros [Kiayias-Russel-David-Oliynykov'17]
- Snow white [Daian-Pass-Shi'17]

Proofs of Space

- Require non-trivial disk space to solve a puzzle
[Dziembowski et al. CRYPTO'15, Ateniese et al. SCN'14]
- More environmental-friendly than PoW
- Used in Chia, and FileCoin (combination of proof of storage/retrievability and proof of space)

Questions with Virtual Mining

Is there any security that can only be gained by consuming “real” resources?

- If so, then “waste” is the cost of security
- If not, then PoW mining may go extinct

Conclusion

- Many possible design goals for puzzles
 - Prevent ASIC miners from dominating
 - Prevent large pools from dominating
 - Intrinsic usefulness
 - Eliminate the need for mining hardware at all
- Further research required to understand the best tradeoffs
- Many competing systems already co-exist