

# Blockchains & Cryptocurrencies

## Proof of Stake - II



Image from cedricwallburger.com

---

Instructor: Abhishek Jain  
Johns Hopkins University - Spring 2021

# Recap: Proof of Stake

- Participants must have some “stake” (i.e., money) in the system
- Chance of “winning” in a block mining round proportional to one’s current stake
- Security requires majority ownership of stake to be honest

# Recap: Agenda

- Algorand
  - Fast Byzantine Agreement (with player replaceability)
  - Private sampling to elect committees
- Other attacks and defenses for PoS systems

# Recap: Byzantine Agreement

- Consider  $n$  parties that have inputs  $\mathbf{v}_i$
- Let  $t$  be the number of maximum corrupted parties
- Communication model: P2P (assume full-connectivity; synchronous)
- Goal: Design an interactive protocol that terminates (with high probability), where
  - Agreement: All honest players output the **same** value
  - Consistency: if all honest players started with same input  $\mathbf{v}$ , then output of all honest players must be  $\mathbf{v}$

# Micali's Protocol: Main Intuition

Consider “idealized” protocol  $\mathbf{P(r)}$ , where  $\mathbf{b_i}$  is the initial input of party  $\mathbf{i}$ :

- Each player  $\mathbf{i}$  sends  $\mathbf{b_i}$  to all other players
- A new random and independently selected bit  $\mathbf{c(r)}$  *appears in sky*
- Player  $\mathbf{i}$  updates bit  $\mathbf{b_i}$  as follows:
  - If  $\#_{i,r}(0) \geq 2t+1$ , set  $\mathbf{b_i} = 0$
  - If  $\#_{i,r}(1) \geq 2t+1$ , set  $\mathbf{b_i} = 1$
  - Else, set  $\mathbf{b_i} = \mathbf{c(r)}$

$\#_{i,r}(\mathbf{b})$ : Number of players from which  $\mathbf{i}$  received  $\mathbf{b}$  in “iteration” number  $\mathbf{r}$

# Quick Analysis

Assume at least  $2t+1$  players are honest:

- If honest players start in agreement, then they remain in agreement
- If honest players do not start in agreement, then they end in agreement (on some bit) with probability  $1/2$
- Think: Why?

# Implementing “coin in sky” using cryptography

## Three Ingredients:

- Unique Digital Signatures: For every public key  $pk$  and message  $m$ , **only one** valid signature for  $m$  w.r.t.  $pk$
- Hash function: Modeled as a random oracle
- Common random string  $R$ : Fixed at the start of the protocol execution, known to each party, and not controlled by adversary

# Implementing “coin in sky” using cryptography

**ConcreteCoin(r)**: Each player  $i$  does the following,

- Send  $s_i = \text{SIG}_i(R, r)$
- Compute  $m$  s.t.  $H(s_m) \leq H(s_i)$  for all  $i$
- Set  $c_i(r) = \text{lsb}(h)$ , where  $h = H(s_m)$

Think: What is the probability that  $c_i(r) = c_j(r)$  for all honest  $i, j$  ?

Think: Why is  $c_i(r)$  random?



# Using ConcreteCoin(r)

Replacing coin in sky with ConcreteCoin(r) in P(r):

- If honest players start in agreement, then they remain in agreement
- If honest players do not start in agreement, then they end in agreement (on some bit) with probability **1/3**

# Remaining Problem

Can we simply repeat the protocol indefinitely until agreement is reached?

- The honest players do not know that agreement is reached
- Thus, the protocol may never terminate

# Remaining Problem

Can we simply repeat the protocol indefinitely until agreement is reached?

- The honest players do not know that agreement is reached
- Thus, the protocol may never terminate

Idea: Simply repeat say  **$k = 200$**  times to ensure that agreement is reached, except with very small probability

Drawback: We waste many rounds since most times, agreement will be reached earlier

# Micali's Idea:

**Protocol BBA\*** : It consists of **sequential repetitions** of  $P'(r)$ , where each  $P'(r)$  consists of three correlated executions of  $P(r)$

1. Execution of  $P(r)$  where  $c(r) = 0$
2. Execution of  $P(r)$  where  $c(r) = 1$
3. Execution of  $P(r)$  where  $c(r)$  is implemented via **ConcreteCoin**( $r$ )

**Note 1:** In the first two executions, a party will **terminate** if it thinks agreement is reached

**Note 2:** While the number of repetitions of  $P'(r)$  are not fixed in advanced, the expected number of repetitions will be 3 (will follow from protocol analysis)

# Notation:

1. A party  $i$  may at any point send special value  $\mathbf{b}^*$  (and HALT) meaning that in all future steps, other parties should think of  $i$ 's message as  $\mathbf{b}$
2. Counter  $\gamma$  which indicates how many times the 3-step loop has been executed. Initially set to  $\mathbf{0}$
3.  $\mathbf{R}$  denotes the common random string

# PROTOCOL $BBA^*$

(COMMUNICATION) STEP 1. [Coin-Fixed-To-0 Step] *Each player  $i$  propagates  $b_i$ .*

1.1 *If  $\#_i^1(0) \geq 2t + 1$ , then  $i$  sets  $b_i = 0$ , sends  $0*$ , outputs  $out_i = 0$ , and HALTS.*

1.2 *If  $\#_i^1(1) \geq 2t + 1$ , then, then  $i$  sets  $b_i = 1$ .*

1.3 *Else,  $i$  sets  $b_i = 0$ .*

(COMMUNICATION) STEP 2. [Coin-Fixed-To-1 Step] *Each player  $i$  propagates  $b_i$ .*

2.1 *If  $\#_i^2(1) \geq 2t + 1$ , then  $i$  sets  $b_i = 1$ , sends  $1*$ , outputs  $out_i = 1$ , and HALTS.*

2.2 *If  $\#_i^2(0) \geq 2t + 1$ , then  $i$  set  $b_i = 0$ .*

2.3 *Else,  $i$  sets  $b_i = 1$ .*

(COMMUNICATION) STEP 3. [Coin-Genuinely-Flipped Step] *Each player  $i$  propagates  $b_i$  and  $SIG_i(R, \gamma)$ .*

3.1 *If  $\#_i^3(0) \geq 2t + 1$ , then  $i$  sets  $b_i = 0$ .*

3.2 *Else, if  $\#_i^3(1) \geq 2t + 1$ , then  $i$  sets  $b_i = 1$ .*

3.3 *Else, letting  $S_i = \{j \in N : m_i^3(j) = SIG_j(R, \gamma)\}$ ,*

*$i$  sets  $b_i = c_i^{(\gamma)} \triangleq \text{lsb}(\min_{j \in S_i} H(SIG_i(R, \gamma)))$ ; increases  $\gamma_i$  by 1; and returns to Step 1.*

# Analysis

**Claim A:** *If at start of an execution of step 3, no player has halted and agreement has not been reached, then with prob  $1/3$ , players will be in agreement at the end of the step*

# Analysis

**Claim A:** If at start of an execution of **step 3**, no player has halted and agreement has not been reached, then with prob **1/3**, players will be in agreement at the end of the step

**Proof:** Consider 5 possible cases:

1. All honest  $i$  update  $b_i$  according to **3.1**
2. All honest  $i$  update  $b_i$  according to **3.2**
3. All honest  $i$  update  $b_i$  according to **3.3**
4. Some honest  $i$  update  $b_i$  according to **3.1**, others according to **3.3**
5. Some honest  $i$  update  $b_i$  according to **3.2**, others according to **3.3**

## PROTOCOL $BBA^*$

(COMMUNICATION) STEP 1. [Coin-Fixed-To-0 Step] *Each player  $i$  propagates  $b_i$ .*

1.1 *If  $\#_i^1(0) \geq 2t + 1$ , then  $i$  sets  $b_i = 0$ , sends  $0^*$ , outputs  $out_i = 0$ , and HALTS.*

1.2 *If  $\#_i^1(1) \geq 2t + 1$ , then, then  $i$  sets  $b_i = 1$ .*

1.3 *Else,  $i$  sets  $b_i = 0$ .*

(COMMUNICATION) STEP 2. [Coin-Fixed-To-1 Step] *Each player  $i$  propagates  $b_i$ .*

2.1 *If  $\#_i^2(1) \geq 2t + 1$ , then  $i$  sets  $b_i = 1$ , sends  $1^*$ , outputs  $out_i = 1$ , and HALTS.*

2.2 *If  $\#_i^2(0) \geq 2t + 1$ , then  $i$  set  $b_i = 0$ .*

2.3 *Else,  $i$  sets  $b_i = 1$ .*

(COMMUNICATION) STEP 3. [Coin-Genuinely-Flipped Step] *Each player  $i$  propagates  $b_i$  and  $SIG_i(R, \gamma)$ .*

3.1 *If  $\#_i^3(0) \geq 2t + 1$ , then  $i$  sets  $b_i = 0$ .*

3.2 *Else, if  $\#_i^3(1) \geq 2t + 1$ , then  $i$  sets  $b_i = 1$ .*

3.3 *Else, letting  $S_i = \{j \in N : m_i^3(j) = SIG_j(R, \gamma)\}$ ,*

*$i$  sets  $b_i = c_i^{(\gamma)} \triangleq \text{lsb}(\min_{j \in S_i} H(SIG_j(R, \gamma)))$ ; increases  $\gamma_i$  by 1; and returns to Step 1.*



**Proof:** Consider 5 possible cases:

1. All honest  $i$  update  $\mathbf{b}_i$  according to **3.1**
  - Agreement holds on **0**
2. All honest  $i$  update  $\mathbf{b}_i$  according to **3.2**
  - Agreement holds on **1**
3. All honest  $i$  update  $\mathbf{b}_i$  according to **3.3**
  - Agreement holds on **c**
4. Some honest  $i$  update  $\mathbf{b}_i$  according to **3.1**, others according to **3.3**
  - Agreement on **0** reached with prob  $1/2$  (assuming  $\mathbf{c}_i$ 's are same)
5. Some honest  $i$  update  $\mathbf{b}_i$  according to **3.2**, others according to **3.3**
  - Agreement on **1** reached with prob  $1/2$  (assuming  $\mathbf{c}_i$ 's are same)

Overall, when  $\mathbf{m}$  is **honest**, agreement is reached with probability at least **1/2** since  $\mathbf{c}_i$ 's are same in this case.  $\mathbf{m}$  is honest with prob **2/3**, so overall agreement prob is **1/3**

(COMMUNICATION) STEP 1. [Coin-Fixed-To-0 Step] Each player  $i$  propagates  $b_i$ .

1.1 If  $\#_i^1(0) \geq 2t + 1$ , then  $i$  sets  $b_i = 0$ , sends  $0^*$ , outputs  $out_i = 0$ , and HALTS.

1.2 If  $\#_i^1(1) \geq 2t + 1$ , then, then  $i$  sets  $b_i = 1$ .

1.3 Else,  $i$  sets  $b_i = 0$ .

(COMMUNICATION) STEP 2. [Coin-Fixed-To-1 Step] Each player  $i$  propagates  $b_i$ .

2.1 If  $\#_i^2(1) \geq 2t + 1$ , then  $i$  sets  $b_i = 1$ , sends  $1^*$ , outputs  $out_i = 1$ , and HALTS.

2.2 If  $\#_i^2(0) \geq 2t + 1$ , then  $i$  set  $b_i = 0$ .

2.3 Else,  $i$  sets  $b_i = 1$ .

(COMMUNICATION) STEP 3. [Coin-Genuinely-Flipped Step] Each player  $i$  propagates  $b_i$  and  $SIG_i(R, \gamma)$ .

3.1 If  $\#_i^3(0) \geq 2t + 1$ , then  $i$  sets  $b_i = 0$ .

3.2 Else, if  $\#_i^3(1) \geq 2t + 1$ , then  $i$  sets  $b_i = 1$ .

3.3 Else, letting  $S_i = \{j \in N : m_i^3(j) = SIG_j(R, \gamma)\}$ ,  
 $i$  sets  $b_i = c_i^{(\gamma)} \triangleq \mathbf{1sb}(\min_{j \in S_i} H(SIG_j(R, \gamma)))$ ; increases  $\gamma_i$  by 1; and returns to Step 1.

## Analysis (contd.)

**Claim B:** *If at some step, agreement holds on bit  $b$ , then it continues to hold on bit  $b$*

## Analysis (contd.)

**Claim B:** *If at some step, agreement holds on bit  $b$ , then it continues to hold on bit  $b$*

**Proof:** If agreement held at the start of step, then all honest parties send bit  $b$ , which means  $\#_i(b) \geq 2t+1$

## Analysis (contd.)

**Claim B:** *If at some step, agreement holds on bit  $b$ , then it continues to hold on bit  $b$*

**Proof:** If agreement held at the start of step, then all honest parties send bit  $b$ , which means  $\#_i(b) \geq 2t+1$

**Claim C:** *If at some step, a player  $i$  halts, then agreement will hold at the end of the step*

# Analysis (contd.)

**Claim B:** *If at some step, agreement holds on bit  $b$ , then it continues to hold on bit  $b$*

**Proof:** If agreement held at the start of step, then all honest parties send bit  $b$ , which means  $\#_i(b) \geq 2t+1$

**Claim C:** *If at some step, a player  $i$  halts, then agreement will hold at the end of the step*

**Proof:** WLOG, suppose  $i$  halts in **Coin-Fixed-To-0** step. Since  $\#_i(0) \geq 2t+1$ , at least  $t+1$  honest players sent  $0$ . Thus,  $\#_j(0) \geq t+1$  for every other honest  $j$ . If  $\#_j(0) \geq 2t+1$ , then  $j$  sets  $b_j=0$  in step 1.1, else it sets  $b_j=0$  in step 1.3. (Step 1.2 cannot be executed)

# Analysis (contd.)

**Property 1:** Consistency (if initial bit  $b$  for all honest players, then  $\text{out}_i = b$ )

**Proof:** Easily follows from step 1.1 or 2.2 (depending upon whether starting input was 0 or 1)

**Property 2:** Agreement ( $\text{out}_i = \text{out}_j$  for all honest  $i, j$ )

**Proof:** Follows from Claims A, B and C

# Player Replaceability

- Parties don't have any protocol “state” across different rounds
- Hence, each round can be executed by different sets of parties
  - Parties “listen” to network
  - When a previous round is completed, check if selected (using private sampling) and participate

Algorand using Byzantine Agreement



# Main Ideas

- Builds on BA protocol of Micali for consensus
- BA protocol executed between a small committee of users for scalability
- Committee chosen at random, using **cryptographic sortition** (aka, private sampling)
- Committee members change in every round of BA protocol

# Main ideas (contd.)

- For every block generation round, a small committee of parties is selected at random based on user stakes
- Each selected party gets to propose a block with some associated “priority”
- Each party distributes its block together with “proof” of selection
- Parties prioritize the blocks based on proposer’s “priority”
- However, different parties may have different views
- To reach consensus on same block, the committee runs BA protocol

# Verifiable Random Functions

- On any input  $\mathbf{x}$ ,  $\text{VRF}_{\text{sk}}(\mathbf{x})$  outputs  $(\text{hash}, \text{proof})$
- **hash** is uniquely determined given **sk** and  $\mathbf{x}$  but indistinguishable from random to anyone who does not know **sk**
- Given **pk** and proof, *anyone can check* that hash corresponds to  $\mathbf{x}$
- Can be built from standard cryptographic assumptions

# Notation

- **W**: total amount of currency units
- **t**: threshold, denoting expected number of users selected
- **p**:  $t/W$
- $w_i$ : stake/money of user **i**
- **B(k;w,p)**: Prob of getting **k** successes in **w** trials, where prob of success in each trial is **p** (Binomial distribution)

$$\sum_{k=0}^w B(k; w, p) = 1$$

- Division of interval **[0,1)** into multiple consecutive intervals

$$I_j = \left[ \sum_{k=0}^j B(k; w, p), \sum_{k=0}^{j+1} B(k; w, p) \right)$$

# Cryptographic Sortition

Sortition(sk,seed,p,w):

- $\text{VRF}_{\text{sk}}(\text{seed}) \rightarrow (\text{hash}, \text{proof})$
- $j \rightarrow 0$
- While  $\frac{\text{hash}}{2^{\text{hashlen}}} \notin \left[ \sum_{k=0}^j B(k; w, p), \sum_{k=0}^{j+1} B(k; w, p) \right)$   
   $j++$
- Return (hash, proof, j)

# Cryptographic Sortition (contd.)

- Any party can compute on its own whether it was selected, together with **proof**, using its **sk**
- A user  $i$  with  $w_i$  units of Algorand is viewed as  $w_i$  potential “sub-users”. Each sub-user is selected with probability  $p=t/W$ . Counter  $j$  denotes how many selected.
- The value of hash determines priority

# Consensus with player replaceability

- For each round of the BA protocol, a fresh set of users is chosen, also using cryptographic sortition
- All users can passively participate in the protocol by listening to the gossip network. Whenever selected for a round, they send a message based on what they heard so far on the network
- BA protocol has player replaceability; therefore using different users in each step is possible

# Security Challenges

- For BA consensus, high majority of players must be honest
- Why can't adversary simply corrupt all the committee members?
- **Main Idea:** Committee members for any step are disclosed only when they send their respective messages. If adversary corrupts now, its too late. The messages are already sent.



# Security Challenges (contd.)

- How to select the threshold  $t$ ?
- Use a threshold such that:
  - $\#good > \text{threshold}$ : for agreement
  - $\frac{1}{2} \#good + \#bad \leq \text{threshold}$ : to avoid forks

## Other Points:

- The seed (used in sortition) has to be chosen carefully. Initially, it is set to be a common random string; later, for each round  $r$ , seed is determined from seed for round  $r-1$  by using  $\mathbf{VRF}_{sk}$  of the block proposer in round  $r-1$
- What are the chances of forks? – Forks can happen with some probability (if network has weak synchrony), but a recovery process can be used to eliminate fork assuming there is a strong synchrony period, using same BA procedure

# Algorand: Summary

- High throughput:  $\sim 1$  min to confirm transactions vs an hour in Bitcoin
- Public ledger with low probability of forks
- Assumes  $2/3$ -honest stake majority
- Uses a gossip communication protocol